


Fine-tuning OpenAI models with Weights & Biases



Anish Shah

Oct 3, 2023

 Open in Github

Note: you will need an [OpenAI API key](#) to run this colab.

If you use OpenAI's API to [fine-tune ChatGPT-3.5](#), you can now use the W&B integration to track experiments, models, and datasets in your central dashboard.

All it takes is one line: `openai wandb sync`

See the [OpenAI section](#) in the Weights & Biases documentation for full details of the integration

```
!pip install -Uq openai tiktoken datasets tenacity wandb
```

```
# Remove once this PR is merged: https://github.com/openai/openai-python/pull/590 and openai release i
!pip uninstall -y openai -qq \
&& pip install git+https://github.com/morganmcg1/openai-python.git@update_wandb_logger -qqq
```

Optional: Fine-tune ChatGPT-3.5

It's always more fun to experiment with your own projects so if you have already used the openai API to fine-tune an OpenAI model, just skip this section.

Otherwise let's fine-tune ChatGPT-3.5 on a legal dataset!

Imports and initial set-up

```
import openai
import wandb

import os
```

```
import json
import random
import tiktoken
import numpy as np
import pandas as pd
from pathlib import Path
from tqdm.auto import tqdm
from collections import defaultdict
from tenacity import retry, stop_after_attempt, wait_fixed
```

Start your Weights & Biases run. If you don't have an account you can sign up for one for free at www.wandb.ai

```
WANDB_PROJECT = "OpenAI-Fine-Tune"
```

Set up your API key

```
# # Enter credentials
openai_key = "YOUR_API_KEY"

openai.api_key = openai_key
```

Dataset Preparation

We download a dataset from [LegalBench](#), a project to curate tasks for evaluating legal reasoning, specifically the [Contract NLI Explicit Identification task](#).

This comprises of a total of 117 examples, from which we will create our own train and test datasets

```
from datasets import load_dataset

# Download the data, merge into a single dataset and shuffle
dataset = load_dataset("nguha/legalbench", "contract_nli_explicit_identification")

data = []
for d in dataset["train"]:
    data.append(d)

for d in dataset["test"]:
    data.append(d)

random.shuffle(data)

for idx, d in enumerate(data):
    d["new_index"] = idx
```

Let's look at a few samples.

```
len(data), data[0:2]
```

```
(117,
 [{ 'answer': 'No',
    'index': '94',
    'text': 'Recipient shall use the Confidential Information exclusively for HySafe purposes,
    'document_name': 'NDA_V3.pdf',
    'new_index': 0},
  { 'answer': 'No',
    'index': '53',
    'text': '3. In consideration of each and every disclosure of CONFIDENTIAL INFORMATION, the
    'document_name': '1084000_0001144204-06-046785_v056501_ex10-16.txt',
    'new_index': 1}])
```

Format our Data for Chat Completion Models

We modify the `base_prompt` from the LegalBench task to make it a zero-shot prompt, as we are training the model instead of using few-shot prompting

```
base_prompt_zero_shot = "Identify if the clause provides that all Confidential Information shall be ex
```

We now split it into training/validation dataset, let's train on 30 samples and test on the remainder

```
n_train = 30
n_test = len(data) - n_train
```

```
train_messages = []
test_messages = []

for d in data:
    prompts = []
    prompts.append({"role": "system", "content": base_prompt_zero_shot})
    prompts.append({"role": "user", "content": d["text"]})
    prompts.append({"role": "assistant", "content": d["answer"]})

    if int(d["new_index"]) < n_train:
        train_messages.append({'messages': prompts})
    else:
        test_messages.append({'messages': prompts})
```

```
len(train_messages), len(test_messages), n_test, train_messages[5]
```

```
(30,
 87,
 87,
 {'messages': [{'role': 'system',
  'content': 'Identify if the clause provides that all Confidential Information shall be ex
 {'role': 'user',
  'content': '2. The Contractor shall not, without the State’s prior written consent, copy,
 {'role': 'assistant', 'content': 'No'}]}}
```

Save the data to Weights & Biases

Save the data in a train and test file first

```
train_file_path = 'encoded_train_data.jsonl'
with open(train_file_path, 'w') as file:
    for item in train_messages:
        line = json.dumps(item)
        file.write(line + '\n')

test_file_path = 'encoded_test_data.jsonl'
with open(test_file_path, 'w') as file:
    for item in test_messages:
        line = json.dumps(item)
        file.write(line + '\n')
```

Next, we validate that our training data is in the correct format using a script from the [OpenAI fine-tuning documentation](#)

```
# Next, we specify the data path and open the JSONL file

def openai_validate_data(dataset_path):
    data_path = dataset_path

    # Load dataset
    with open(data_path) as f:
        dataset = [json.loads(line) for line in f]

    # We can inspect the data quickly by checking the number of examples and the first item

    # Initial dataset stats
    print("Num examples:", len(dataset))
    print("First example:")
    for message in dataset[0]["messages"]:
        print(message)

    # Now that we have a sense of the data, we need to go through all the different examples and check t
```

```

# Format error checks
format_errors = defaultdict(int)

for ex in dataset:
    if not isinstance(ex, dict):
        format_errors["data_type"] += 1
        continue

    messages = ex.get("messages", None)
    if not messages:
        format_errors["missing_messages_list"] += 1
        continue

    for message in messages:
        if "role" not in message or "content" not in message:
            format_errors["message_missing_key"] += 1

        if any(k not in ("role", "content", "name") for k in message):
            format_errors["message_unrecognized_key"] += 1

        if message.get("role", None) not in ("system", "user", "assistant"):
            format_errors["unrecognized_role"] += 1

        content = message.get("content", None)
        if not content or not isinstance(content, str):
            format_errors["missing_content"] += 1

    if not any(message.get("role", None) == "assistant" for message in messages):
        format_errors["example_missing_assistant_message"] += 1

if format_errors:
    print("Found errors:")
    for k, v in format_errors.items():
        print(f"{k}: {v}")
else:
    print("No errors found")

# Beyond the structure of the message, we also need to ensure that the length does not exceed the 40

# Token counting functions
encoding = tiktoken.get_encoding("cl100k_base")

# not exact!
# simplified from https://github.com/openai/openai-cookbook/blob/main/examples/How_to_count_tokens_w
def num_tokens_from_messages(messages, tokens_per_message=3, tokens_per_name=1):
    num_tokens = 0
    for message in messages:
        num_tokens += tokens_per_message
        for key, value in message.items():
            num_tokens += len(encoding.encode(value))
            if key == "name":
                num_tokens += tokens_per_name
    num_tokens += 3
    return num_tokens

def num_assistant_tokens_from_messages(messages):
    num_tokens = 0
    for message in messages:
        if message["role"] == "assistant":
            num_tokens += len(encoding.encode(message["content"]))
    return num_tokens

```

```

def print_distribution(values, name):
    print(f"\n#### Distribution of {name}:")
    print(f"min / max: {min(values)}, {max(values)}")
    print(f"mean / median: {np.mean(values)}, {np.median(values)}")
    print(f"p5 / p95: {np.quantile(values, 0.1)}, {np.quantile(values, 0.9)}")

# Last, we can look at the results of the different formatting operations before proceeding with cre

# Warnings and tokens counts
n_missing_system = 0
n_missing_user = 0
n_messages = []
convo_lens = []
assistant_message_lens = []

for ex in dataset:
    messages = ex["messages"]
    if not any(message["role"] == "system" for message in messages):
        n_missing_system += 1
    if not any(message["role"] == "user" for message in messages):
        n_missing_user += 1
    n_messages.append(len(messages))
    convo_lens.append(num_tokens_from_messages(messages))
    assistant_message_lens.append(num_assistant_tokens_from_messages(messages))

print("Num examples missing system message:", n_missing_system)
print("Num examples missing user message:", n_missing_user)
print_distribution(n_messages, "num_messages_per_example")
print_distribution(convo_lens, "num_total_tokens_per_example")
print_distribution(assistant_message_lens, "num_assistant_tokens_per_example")
n_too_long = sum(l > 4096 for l in convo_lens)
print(f"\n{n_too_long} examples may be over the 4096 token limit, they will be truncated during fine

# Pricing and default n_epochs estimate
MAX_TOKENS_PER_EXAMPLE = 4096

MIN_TARGET_EXAMPLES = 100
MAX_TARGET_EXAMPLES = 25000
TARGET_EPOCHS = 3
MIN_EPOCHS = 1
MAX_EPOCHS = 25

n_epochs = TARGET_EPOCHS
n_train_examples = len(dataset)
if n_train_examples * TARGET_EPOCHS < MIN_TARGET_EXAMPLES:
    n_epochs = min(MAX_EPOCHS, MIN_TARGET_EXAMPLES // n_train_examples)
elif n_train_examples * TARGET_EPOCHS > MAX_TARGET_EXAMPLES:
    n_epochs = max(MIN_EPOCHS, MAX_TARGET_EXAMPLES // n_train_examples)

n_billing_tokens_in_dataset = sum(min(MAX_TOKENS_PER_EXAMPLE, length) for length in convo_lens)
print(f"Dataset has ~{n_billing_tokens_in_dataset} tokens that will be charged for during training")
print(f"By default, you'll train for {n_epochs} epochs on this dataset")
print(f"By default, you'll be charged for ~{n_epochs * n_billing_tokens_in_dataset} tokens")
print("See pricing page to estimate total costs")

```

Validate train data

```
openai_validate_data(train_file_path)
```

```
Num examples: 30
First example:
{'role': 'system', 'content': 'Identify if the clause provides that all Confidential Information shall be used exclusively for the purposes of the business of the recipient.'}
{'role': 'user', 'content': 'Recipient shall use the Confidential Information exclusively for the purposes of the business of the recipient.'}
{'role': 'assistant', 'content': 'No'}
No errors found
Num examples missing system message: 0
Num examples missing user message: 0

#### Distribution of num_messages_per_example:
min / max: 3, 3
mean / median: 3.0, 3.0
p5 / p95: 3.0, 3.0

#### Distribution of num_total_tokens_per_example:
min / max: 69, 319
mean / median: 143.46666666666667, 122.0
p5 / p95: 82.10000000000001, 235.10000000000002

#### Distribution of num_assistant_tokens_per_example:
min / max: 1, 1
mean / median: 1.0, 1.0
p5 / p95: 1.0, 1.0

0 examples may be over the 4096 token limit, they will be truncated during fine-tuning
Dataset has ~4304 tokens that will be charged for during training
By default, you'll train for 3 epochs on this dataset
By default, you'll be charged for ~12912 tokens
See pricing page to estimate total costs
```

Log our data to Weights & Biases Artifacts for storage and versioning

```
wandb.init(
    project=WANDB_PROJECT,
    # entity="prompt-eng",
    job_type="log-data",
    config = {'n_train': n_train,
              'n_valid': n_test})

wandb.log_artifact(train_file_path,
                   "legalbench-contract_nli_explicit_identification-train",
                   type="train-data")

wandb.log_artifact(test_file_path,
                   "legalbench-contract_nli_explicit_identification-test",
                   type="test-data")

# keep entity (typically your wandb username) for reference of artifact later in this demo
entity = wandb.run.entity

wandb.finish()
```

Failed to detect the name of this notebook, you can set it manually with the WANDB_NOTEBOOK_NAME

```
[34m [1mwandb [0m: Currently logged in as: [33mcapecape [0m. Use [1m`wandb login --relogin` [0
```

Tracking run with wandb version 0.15.9

Run data is saved locally in `/Users/tcapelle/work/examples/colabs/openai/wandb/run-20230830_113853-ivu21mjl`

Syncing run `mild-surf-1` to [Weights & Biases \(docs\)](#)

View project at <https://wandb.ai/capecape/OpenAI-Fine-Tune>

View run at <https://wandb.ai/capecape/OpenAI-Fine-Tune/runs/ivu21mjl>

Create a fine-tuned model

We'll now use OpenAI API to fine-tune ChatGPT-3.5

Let's first download our training & validation files and save them to a folder called `my_data`. We will retrieve the `latest` version of the artifact, but it could also be `v0`, `v1` or any alias we associated with it

```
wandb.init(project=WANDB_PROJECT,
           # entity="prompt-eng",
           job_type="finetune")

artifact_train = wandb.use_artifact(
    f'{entity}/{WANDB_PROJECT}/legalbench-contract_nli_explicit_identification-train:latest',
    type='train-data')
train_file = artifact_train.get_path(train_file_path).download("my_data")

train_file
```

```
VBox(children=(Label(value='Waiting for wandb.init()...\r'), FloatProgress(value=0.01675180270
```


Tracking run with wandb version 0.15.9

Run data is saved locally in `/Users/tcapelle/work/examples/colabs/openai/wandb/run-20230830_113907-1ili9l51`

Syncing run `jumping-water-2` to [Weights & Biases \(docs\)](#)

View project at <https://wandb.ai/capecape/OpenAI-Fine-Tune>

Then we upload the training data to OpenAI. OpenAI has to process the data, so this will take a few minutes depending on the size of your dataset.

```
openai_train_file_info = openai.File.create(
    file=open(train_file, "rb"),
    purpose='fine-tune'
)

# you may need to wait a couple of minutes for OpenAI to process the file
openai_train_file_info
```

```
<File file id=file-spPASR6VWco54SqfN2yo7T8v> JSON: {
  "object": "file",
  "id": "file-spPASR6VWco54SqfN2yo7T8v",
  "purpose": "fine-tune",
  "filename": "file",
  "bytes": 24059,
  "created_at": 1693388388,
  "status": "uploaded",
  "status_details": null
}
```

Time to train the model!

Let's define our ChatGPT-3.5 fine-tuning hyper-parameters.

```
model = 'gpt-3.5-turbo'
n_epochs = 3
```



```
openai_ft_job_info = openai.FineTuningJob.create(
    training_file=openai_train_file_info["id"],
    model=model,
    hyperparameters={"n_epochs": n_epochs}
)

ft_job_id = openai_ft_job_info["id"]

openai_ft_job_info
```



```
<FineTuningJob fine_tuning.job id=ftjob-x4tl83IlsGolkUF3fCFyZNGs> JSON: {
  "object": "fine_tuning.job",
  "id": "ftjob-x4tl83IlsGolkUF3fCFyZNGs",
  "model": "gpt-3.5-turbo-0613",
  "created_at": 1693388447,
  "finished_at": null,
  "fine_tuned_model": null,
  "organization_id": "org-WnF2wEqNkV1Nj65CzDxr6iUm",
  "result_files": [],
  "status": "created",
  "validation_file": null,
  "training_file": "file-spPASR6VWco54SqfN2yo7T8v",
  "hyperparameters": {
    "n_epochs": 3
  },
  "trained_tokens": null
}
```

"this takes around 5 minutes to train, and you get an email from OpenAI when finished."

Thats it!

Now your model is training on OpenAI's machines. To get the current state of your fine-tuning job, run:

```
state = openai.FineTuningJob.retrieve(ft_job_id)
state["status"], state["trained_tokens"], state["finished_at"], state["fine_tuned_model"]
```



```
('succeeded',
 12732,
```

```
1693389024,  
'ft:gpt-3.5-turbo-0613:weights-biases::7tC85HcX')
```

Show recent events for our fine-tuning job

```
openai.FineTuningJob.list_events(id=ft_job_id, limit=5)
```

```
<OpenAIObject list> JSON: {  
  "object": "list",  
  "data": [  
    {  
      "object": "fine_tuning.job.event",  
      "id": "ftevent-5x9Y6Payk6fIdyJyMRY5um1v",  
      "created_at": 1693389024,  
      "level": "info",  
      "message": "Fine-tuning job successfully completed",  
      "data": null,  
      "type": "message"  
    },  
    {  
      "object": "fine_tuning.job.event",  
      "id": "ftevent-i16NTGNakv9P0Rk0tJ7vvvoG",  
      "created_at": 1693389022,  
      "level": "info",  
      "message": "New fine-tuned model created: ft:gpt-3.5-turbo-0613:weights-biases::7tC85Hc",  
      "data": null,  
      "type": "message"  
    },  
    {  
      "object": "fine_tuning.job.event",  
      "id": "ftevent-MkLrJQ8sDgaC67CdmFMwsIjV",  
      "created_at": 1693389017,  
      "level": "info",  
      "message": "Step 90/90: training loss=0.00",  
      "data": {  
        "step": 90,  

```

We can run a few different fine-tunes with different parameters or even with different datasets.

Log OpenAI fine-tune jobs to Weights & Biases

We can log our fine-tunes with a simple command.

```
!openai wandb sync --help
```

```
usage: openai wandb sync [-h] [-i ID] [-n N_FINE_TUNES] [--project PROJECT]
                        [--entity ENTITY] [--force] [--legacy]

options:
  -h, --help            show this help message and exit
  -i ID, --id ID        The id of the fine-tune job (optional)
  -n N_FINE_TUNES, --n_fine_tunes N_FINE_TUNES
                        Number of most recent fine-tunes to log when an id is
                        not provided. By default, every fine-tune is synced.
  --project PROJECT     Name of the Weights & Biases project where you're
                        sending runs. By default, it is "OpenAI-Fine-Tune".
  --entity ENTITY       Weights & Biases username or team name where you're
                        sending runs. By default, your default entity is used,
                        which is usually your username.
  --force               Forces logging and overwrite existing wandb run of the
                        same fine-tune.
  --legacy              Log results from legacy OpenAI /v1/fine-tunes api
```

Calling `openai wandb sync` will log all un-synced fine-tuned jobs to W&B

Below we are just logging 1 job, passing:

- our OpenAI key as an environment variable
- the id of the fine-tune job we'd like to log
- the W&B project of where to log it to

See the [OpenAI section](#) in the Weights & Biases documentation for full details of the integration

```
!OPENAI_API_KEY={openai_key} openai wandb sync --id {ft_job_id} --project {WANDB_PROJECT}
```

Retrieving fine-tune job...

```
[34m [1mwandb [0m: Currently logged in as: [33mcapecape [0m. Use [1m`wandb login --relogin` [0
[34m [1mwandb [0m: Tracking run with wandb version 0.15.9
[34m [1mwandb [0m: Run data is saved locally in [35m [1m/Users/tcapelle/work/examples/colabs/c
[34m [1mwandb [0m: Run [1m`wandb offline` [0m to turn off syncing.
[34m [1mwandb [0m: Syncing run [33mftjob-x4tl83ILSGolkUF3fCFyZNGs [0m
[34m [1mwandb [0m: ☆ View project at [34m [4mhttps://wandb.ai/capecape/OpenAI-Fine-Tune [0m
[34m [1mwandb [0m: 🚀 View run at [34m [4mhttps://wandb.ai/capecape/OpenAI-Fine-Tune/runs/ftjc
[34m [1mwandb [0m: Waiting for W&B process to finish... [32m(success). [0m
[34m [1mwandb [0m:
[34m [1mwandb [0m: Run history:
[34m [1mwandb [0m: train_accuracy ████████████████████████████████████████████████████████
[34m [1mwandb [0m:      train_loss ████████████████████████████████████████████████████████
[34m [1mwandb [0m:
[34m [1mwandb [0m: Run summary:
[34m [1mwandb [0m: fine_tuned_model ft:gpt-3.5-turbo-061...
[34m [1mwandb [0m:      status succeeded
[34m [1mwandb [0m:      train_accuracy 1.0
```

```
[34m [1mwandb [0m:      train_loss 0.0
[34m [1mwandb [0m:
[34m [1mwandb [0m: 🚀 View run [33mftjob-x4tl83ILSGolkUF3fCFyZNGs [0m at: [34m [4mhttps://wandb.ai/capecape/OpenAI-Fine-Tune/runs/1ili9l51
[34m [1mwandb [0m: Synced 6 W&B file(s), 0 media file(s), 1 artifact file(s) and 0 other file(s)
[34m [1mwandb [0m: Find logs at: [35m [1m./wandb/run-20230830_115915-ftjob-x4tl83ILSGolkUF3fCFyZNGs
🎉 wandb sync completed successfully
```

```
wandb.finish()
```

Waiting for W&B process to finish... **(success)**.

VBox(children=(Label(value='0.050 MB of 0.050 MB uploaded (0.000 MB deduped)\r'), FloatProgress(value=0.0, max=1.0)),)

```
wandb: WARNING Source type is set to 'repo' but some required information is missing from the
upload_file exception https://storage.googleapis.com/wandb-production.appspot.com/capecape/OpenAI-Fine-Tune/runs/1ili9l51
upload_file request headers: {'User-Agent': 'python-requests/2.28.2', 'Accept-Encoding': 'gzip, deflate', 'Host': 'storage.googleapis.com', 'Authorization': 'Bearer [REDACTED]'}
upload_file response body:
upload_file exception https://storage.googleapis.com/wandb-production.appspot.com/capecape/OpenAI-Fine-Tune/runs/1ili9l51
upload_file request headers: {'User-Agent': 'python-requests/2.28.2', 'Accept-Encoding': 'gzip, deflate', 'Host': 'storage.googleapis.com', 'Authorization': 'Bearer [REDACTED]'}
upload_file response body:
```

View run **jumping-water-2** at: <https://wandb.ai/capecape/OpenAI-Fine-Tune/runs/1ili9l51>
Synced 7 W&B file(s), 0 media file(s), 0 artifact file(s) and 1 other file(s)

Find logs at: `./wandb/run-20230830_113907-1ili9l51/logs`

Our fine-tunes are now successfully synced to Weights & Biases.



Anytime we have new fine-tunes, we can just call `openai wandb sync` to add them to our dashboard.

Run evaluation and log the results

The best way to evaluate a generative model is to explore sample predictions from your evaluation set.

Let's generate a few inference samples and log them to W&B and see how the performance compares to a baseline ChatGPT-3.5 model

```
wandb.init(project=WANDB_PROJECT,
           job_type='eval')

artifact_valid = wandb.use_artifact(
    f'{entity}/{WANDB_PROJECT}/legalbench-contract_nli_explicit_identification-test:latest',
    type='test-data')
test_file = artifact_valid.get_path(test_file_path).download("my_data")

with open(test_file) as f:
    test_dataset = [json.loads(line) for line in f]

print(f"There are {len(test_dataset)} test examples")
wandb.config.update({"num_test_samples": len(test_dataset)})
```

Tracking run with wandb version 0.15.9

Run data is saved locally in `/Users/tcapelle/work/examples/colabs/openai/wandb/run-20230830_115947-
iepk19m2`

Syncing run `ethereal-energy-4` to [Weights & Biases](#) ([docs](#))

View project at <https://wandb.ai/capecape/OpenAI-Fine-Tune>

View run at <https://wandb.ai/capecape/OpenAI-Fine-Tune/runs/iepk19m2>

There are 87 test examples

Run evaluation on the Fine-Tuned Model

Set up OpenAI call with retries

```
@retry(stop=stop_after_attempt(3), wait=wait_fixed(60))
def call_openai(messages="", model="gpt-3.5-turbo"):
    return openai.ChatCompletion.create(model=model, messages=messages, max_tokens=10)
```

Let's get our trained model id

```
state = openai.FineTuningJob.retrieve(ft_job_id)
ft_model_id = state["fine_tuned_model"]
ft_model_id
```

```
'ft:gpt-3.5-turbo-0613:weights-biases::7tC85HcX'
```

Run evaluation and log results to W&B

```
prediction_table = wandb.Table(columns=['messages', 'completion', 'target'])
eval_data = []

for row in tqdm(test_dataset):
    messages = row['messages'][:2]
    target = row["messages"][2]

    # res = call_openai(model=ft_model_id, messages=messages)
    res = openai.ChatCompletion.create(model=model, messages=messages, max_tokens=10)
    completion = res.choices[0].message.content

    eval_data.append([messages, completion, target])
    prediction_table.add_data(messages[1]['content'], completion, target["content"])

wandb.log({'predictions': prediction_table})
```

```
0%|          | 0/87 [00:00<?, ?it/s]
```

Calculate the accuracy of the fine-tuned model and log to W&B

```
correct = 0
for e in eval_data:
    if e[1].lower() == e[2]["content"].lower():
        correct+=1

accuracy = correct / len(eval_data)
```

```
print(f"Accuracy is {accuracy}")
```

```
Accuracy is 0.8390804597701149
```

Run evaluation on a Baseline model for comparison

Lets compare our model to the baseline model, `gpt-3.5-turbo`

```
baseline_prediction_table = wandb.Table(columns=['messages', 'completion', 'target'])
baseline_eval_data = []

for row in tqdm(test_dataset):
    messages = row['messages'][:2]
    target = row["messages"][2]

    res = call_openai(model="gpt-3.5-turbo", messages=messages)
    completion = res.choices[0].message.content

    baseline_eval_data.append([messages, completion, target])
    baseline_prediction_table.add_data(messages[1]['content'], completion, target["content"])

wandb.log({'baseline_predictions': baseline_prediction_table})
```

```
0%|          | 0/87 [00:00<?, ?it/s]
```

Calculate the accuracy of the fine-tuned model and log to W&B

```
baseline_correct = 0
for e in baseline_eval_data:
    if e[1].lower() == e[2]["content"].lower():
        baseline_correct+=1

baseline_accuracy = baseline_correct / len(baseline_eval_data)
print(f"Baseline Accurcy is: {baseline_accuracy}")
wandb.log({"eval/baseline_accuracy": baseline_accuracy})
wandb.summary["eval/baseline_accuracy"] = baseline_accuracy
```


Baseline Accuracy is: 0.7931034482758621

```
wandb.finish()
```

Waiting for W&B process to finish... (success).

VBox(children=(Label(value='0.248 MB of 0.248 MB uploaded (0.000 MB deduped)\r'), FloatProgress

wandb: WARNING Source type is set to 'repo' but some required information is missing from the

Run history:

eval/accuracy	—
eval/baseline_accuracy	—

Run summary:

eval/accuracy	0.83908
eval/baseline_accuracy	0.7931

And thats it! In this example we have prepared our data, logged it to Weights & Biases, fine-tuned an OpenAI model using that data, logged the results to Weights & Biases and then run evaluation on the fine-tuned model.

From here you can start to train on larger or more complex tasks, or else explore other ways to modify ChatGPT-3.5 such as giving it a different tone and style or response.

Resources

- [OpenAI Fine-Tuning Guide](#)

- [W&B Integration with OpenAI API Documentation](#)
- [W&B Report: GPT-3 exploration & fine-tuning tips](#)